

Hive Assignment

Step 1 : Get data files in Master node

wget <https://e-commerce-events-ml.s3.amazonaws.com/2019-Oct.csv>

wget <https://e-commerce-events-ml.s3.amazonaws.com/2019-Nov.csv>

Step 2 : remove UTC from the timestamp column before loading the data

```
sed -i -e 's/ UTC//g' 2019-Oct.csv
```

```
sed -i -e 's/ UTC//g' 2019-Nov.csv
```

Step 3 : Import csv files in HDFS by running following steps

```
hdfs dfs -mkdir /user/hive_assignment
```

```
hdfs dfs -put 2019-Nov.csv /user/hive_assignment/
```

```
hdfs dfs -put 2019-Oct.csv /user/hive_assignment/
```

Step 4 : Verify if the files are imported in HDFS successfully.

```
hdfs dfs -ls /user/hive_assignment/
```

```
hdfs dfs -cat /user/hive_assignment/2019-Nov.csv | head
```

```
hdfs dfs -cat /user/hive_assignment/2019-Oct.csv | head
```

Step 5 : Enter into Hive console by typing 'hive'

Step 6 : Create database and required tables using following commands

```
create database if not exists ECom comment "Created to store events done on website" with  
dbproperties('creator'='Prakash','date'='30-01-2021');
```

```
show databases;
```

```
use ecom;
```

```
create external table if not exists event_nov (event_time timestamp,event_type string,product_id  
string,category_id string,category_code string,brand string,price float,user_id bigint,user_session  
string) row format delimited fields terminated by ',' lines terminated by '\n' stored as textfile  
tblproperties("skip.header.line.count"="1");
```

```
create external table if not exists event_oct (event_time timestamp,event_type string,product_id  
string,category_id string,category_code string,brand string,price float,user_id bigint,user_session  
string) row format delimited fields terminated by ',' lines terminated by '\n' stored as textfile  
tblproperties("skip.header.line.count"="1");
```

```
show tables;
```

Step 7 : Load data to both tables created for October and November

```
load data inpath '/user/hive_assignment/2019-Nov.csv' into table event_nov;
```

```
load data inpath '/user/hive_assignment/2019-Oct.csv' into table event_oct;
```

```
set hive.cli.print.header=true;
```

```
select * from event_nov limit 10;
```

```
select * from event_nov limit 10;
```

Step 8 : Consolidate the data of October and November in a single table

```
create table events
```

```
as
```

```
(
```

```
select 'Nov' as Month,* from event_nov
```

```
union all
```

```
select 'Oct' as Month,* from event_oct
```

```
)
```

Step 9 : Create partitioning(Static partitioning was used) and bucketing. (I have chosen month and event_type for partitioning and product_id for bucketing. I chose this because most of the queries were based on event type and month. Both columns were having low cardinality)

```
create table if not exists part_events (event_time timestamp,product_id string,category_id string,category_code string,brand string,price float,user_id bigint,user_session string) partitioned by (month string,event_type string) clustered by (product_id) into 5 buckets row format delimited fields terminated by ',' lines terminated by '\n';
```

```
insert into part_events partition( month='Oct', event_type='cart') select event_time,product_id,category_id,category_code,brand,price,user_id, user_session from events where month='Oct' and event_type='cart';
```

```
insert into part_events partition( month='Oct', event_type='remove_from_cart') select event_time,product_id,category_id,category_code,brand,price,user_id, user_session from events where month='Oct' and event_type='remove_from_cart';
```

```
insert into part_events partition( month='Oct',event_type='view') select event_time,product_id,category_id,category_code,brand,price,user_id, user_session from events where month='Oct' and event_type='view';
```

```
insert into part_events partition( month='Oct',event_type='purchase') select
event_time,product_id,category_id,category_code,brand,price,user_id, user_session from events
where month='Oct' and event_type='purchase';
```

```
insert into part_events partition( month='Nov', event_type='cart') select
event_time,product_id,category_id,category_code,brand,price,user_id, user_session from events
where month='Nov' and event_type='cart';
```

```
insert into part_events partition(month='Nov',event_type='remove_from_cart') select
event_time,product_id,category_id,category_code,brand,price,user_id, user_session from events
where month='Nov' and event_type='remove_from_cart';
```

```
insert into part_events partition(month='Nov',event_type='view') select
event_time,product_id,category_id,category_code,brand,price,user_id, user_session from events
where month='Nov' and event_type='view';
```

```
insert into part_events partition(month='Nov',event_type='purchase') select
event_time,product_id,category_id,category_code,brand,price,user_id, user_session from events
where month='Nov' and event_type='purchase';
```

Step 10 : Verify if partitioning and bucketing is done

```
hdfs dfs -ls /user/hive/warehouse/ecom.db/part_events
```

Step 11 : Run a test if the query will be optimized after partitioning.(I have checked the performance for question #1)

I have queried the events table directly which is not partitioned and it is taking 21 seconds. However when I used the partitioned table it is giving result in 12 seconds. There is a improvement of 43 percent. Please find the screenshot below.

`select sum(price) as Total_Revenue_from_purchases from events where month='Oct' and event_type='purchase';`

The screenshot shows the Hive web interface. At the top, there's a header with the Hive logo, a refresh button, and fields for "Add a name..." and "Add a description...". On the right, there are icons for a chart, a document, and a menu. Below the header, the query is displayed: `1 select sum(price) as Total_Revenue_from_purchases from part_events where month='Oct' and event_type='purchase';`. The status bar indicates "20.97s Database ecom Type text". Below the query, the execution log shows: `INFO : Map 1: 4/4 Reducer 2: 1/1`, `INFO : Completed executing command(queryId=hive_20210206090924_4a8211fa-8f25-4626-900d-0844402ffc2f); Time taken: 20.818 seconds`, and `INFO : OK`. A link to the application ID is visible. At the bottom, there are tabs for "Query History", "Saved Queries", "Query Builder", and "Results (1)". The "Results (1)" tab is active, showing a table with one row:

total_revenue_from_purchases	
1	1211538.4295325726

`select sum(price) as Total_Revenue_from_purchases from part_events where month='Oct' and event_type='purchase';`

The screenshot shows the Hive web interface. At the top, there's a header with the Hive logo, a refresh button, and fields for "Add a name..." and "Add a description...". On the right, there are icons for a chart, a document, and a menu. Below the header, the query is displayed: `1 select sum(price) as Total_Revenue_from_purchases from part_events where month='Oct' and event_type='purchase';`. The status bar indicates "12.20s Database ecom Type text". Below the query, the execution log shows: `INFO : Map 1: 3/3 Reducer 2: 1/1`, `INFO : Completed executing command(queryId=hive_20210206091007_2bdca2d2-238c-4499-945e-f7855561a058); Time taken: 11.764 seconds`, and `INFO : OK`. A link to the application ID is visible. At the bottom, there are tabs for "Query History", "Saved Queries", "Query Builder", and "Results (1)". The "Results (1)" tab is active, showing a table with one row:

total_revenue_from_purchases	
1	1211538.4295325726

Step 12 : Answer the Questions asked. Please find below the Answers of all queries/questions.

1. Find the total revenue generated due to the purchases made in October.

```
select sum(price) as Total_Revenue_from_purchases from part_events where month='Oct' and event_type='purchase';
```

The screenshot shows the Hive web interface. At the top, there's a header with the Hive logo, a refresh button, and fields for 'Add a name...' and 'Add a description...'. On the right, there are icons for a chart, a document, and a menu. Below the header, the query is displayed in a text area: `1 select sum(price) as Total_Revenue_from_purchases from part_events where month='Oct' and event_type='purchase';`. The status bar indicates '12.20s Database ecom Type text'. Below the query, there's a log section showing execution progress: 'INFO : Map 1: 3/3 Reducer 2: 1/1', 'INFO : Completed executing command(queryId=hive_20210206091007_2bdca2d2-238c-4499-945e-f78555e1a068); Time taken: 11.764 seconds', and 'INFO : OK'. A link to 'application_1612593054396_0010' is also visible. At the bottom, there are tabs for 'Query History', 'Saved Queries', 'Query Builder', and 'Results (1)'. The 'Results (1)' tab is active, showing a table with the header 'total_revenue_from_purchases' and one row with the value '1211538.4295325726'.

total_revenue_from_purchases
1211538.4295325726

2. Write a query to yield the total sum of purchases per month in a single output.

```
select month, sum(price) as Sum_of_purchases from part_events where event_type='purchase' group by month ;
```

Hive
★
🔄
Add a name...
Add a description...

4.26s Database ecom Type text ⚙️ ?

```

1 select month, sum(price) as Sum_of_purchases from part_events where event_type='purchase' group by month ;

```

```

INFO : Map 1: 5/5 Reducer 2: 1/1/2
INFO : Completed executing command(queryId=hive_20210206092000_16864aae-e338-4f80-8720-6609c3eb7433); Time taken: 3.326 seconds
INFO : OK

```

Query History
Saved Queries
Query Builder
Results (2)

	month	sum_of_purchases
1	Nov	1531016.8991247676
2	Oct	1211538.4295325726

- Write a query to find the change in the revenue generated due to purchases made from October to November.

with MonthlyRevenue as
(select month, sum(price) as Revenue_purchases
from part_events
where event_type='purchase'
group by month)

select (nov.Revenue_purchases-oct.Revenue_purchases) as change_in_revenue
from MonthlyRevenue nov cross join MonthlyRevenue oct
where nov.month<>oct.month and nov.month='Nov' and oct.month='Oct';

Add a name...
Add a description...

7.79s Database ecom Type text ?

```

1 with MonthlyRevenue as
2   (select month, sum(price) as Revenue_purchases
3    from part_events
4    where event_type='purchase'
5    group by month)
6
7 select (nov.Revenue_purchases-oct.Revenue_purchases) as change_in_revenue
8 from MonthlyRevenue nov cross join MonthlyRevenue oct
9 where nov.month<>oct.month and nov.month='Nov' and oct.month='Oct';

```

Map 1: 5/5

Map 3: 3/3

Reducer 2: 2/2

Reducer 4: 2/2

INFO : Completed executing command(queryId=hive_20210206121110_08cb23cd-cb9a-4d53-b7a1-27247b66a742), Time taken:

7.792 seconds

INFO : OK

application_1612593054396_0016

Query History
Saved Queries
Query Builder
Results (1)



change_in_revenue



1	319478.469592195
---	------------------

Note: use set hive.strict.checks.cartesian.product = False for cross product.

- Find distinct categories of products.

```
select distinct category_id from part_events;
```

25.10s Database ecom ▾ Type text ▾  

1

```
select distinct category_id from part_events;
```

INFO : Map 1: 6/6 Reducer 2: 1/1

INFO : Completed executing command(queryId=hive_20210206101510_961f1d27-c500-4f00-803f-3f5d233d3b0c); Time taken: 24.369 seconds

INFO : OK

Query History Saved Queries Query Builder Results (100+)

	category_id
1	1487580004832248652
2	1487580004857414477
3	1487580004882580302
4	1487580004916134735
5	1487580004966466385

5. Find the total number of products available under each category.

```
with categories as (select distinct category_id, product_id from part_events)
select category_id, count(category_id) from categories group by category_id;
```

The screenshot shows the Hive query interface. At the top, there's a header with the Hive logo, a refresh button, and fields for 'Add a name...' and 'Add a description...'. Below this, the query is entered in a text area. The query is:
1 with categories as
2 (select distinct category_id, product_id
3 from part_events)
4
5 select category_id, count(product_id) as number_of_product
6 from categories
7 group by category_id;
8
The query is executed, and the results are displayed in a table. The table has two columns: 'category_id' and 'number_of_product'. The results are as follows:
1 1487580004832248652 283
2 1487580004857414477 543
3 1487580004882580302 148
4 1487580004916134735 585
5 1487580004966466385 9
6 1487580004983243602 43
7 1487580005008409427 460
8 1487580005025186644 1
The interface also shows a log of the query execution, indicating that the query was completed successfully in 12.124 seconds.

```
1 with categories as
2 (select distinct category_id, product_id
3 from part_events)
4
5 select category_id, count(product_id) as number_of_product
6 from categories
7 group by category_id;
8
```

INFO : Map 1: 6/6 Reducer 2: 1/1
INFO : Completed executing command(queryId=hive_20210206122139_e3589b27-e4ba-4275-90c1-ad11d574b054); Time taken: 12.124 seconds
INFO : OK

category_id	number_of_product
1 1487580004832248652	283
2 1487580004857414477	543
3 1487580004882580302	148
4 1487580004916134735	585
5 1487580004966466385	9
6 1487580004983243602	43
7 1487580005008409427	460
8 1487580005025186644	1

6. Which brand had the maximum sales in October and November combined?

```
select brand, sum(price) as revenue from part_events where event_type='purchase' and
brand <> ""
group by brand order by revenue desc limit 1;
```

Add a name...
Add a description...

4.10s
Database ecom
Type text
⚙️
?

```

1 select brand, sum(price) as revenue from part_events where event_type='purchase' and brand <> ""
2 group by brand order by revenue desc limit 1;

```

▶
📖

```

INFO : Map 1: 5/5    Reducer 2: 2/2    Reducer 3: 1/1
INFO : Completed executing command(queryId=hive_20210206111909_ffa6d3fa-a5dc-4b03-a17d-594cb3306c0d), Time taken: 3.266 seconds
INFO : OK

```

application_1612593054396_0015

Query History
Saved Queries
Query Builder
Results (1)

	brand	revenue
1	runail	148297.93996394053

🔍
📄
📁

7. Which brands increased their sales from October to November?

```

with BrandRevenue as
(select month, brand ,sum(price) as sales_per_brand
from part_events
where event_type='purchase'
and brand <> ""
group by month,brand)

select nov.brand
from BrandRevenue nov
inner join BrandRevenue oct
on nov.brand =oct.brand and nov.month <> oct.month
where nov.sales_per_brand >oct.sales_per_brand

```

Add a name...
Add a description...

5.10s Database ecom ▼ Type text ▼ ⚙️ ?

```

1 with BrandRevenue as
2 (select month, brand ,sum(price) as sales_per_brand
3  from part_events
4  where event_type='purchase'
5  and brand <> ""
6  group by month,brand)
7
8 select nov.brand
9  from BrandRevenue nov
10 inner join BrandRevenue oct
11 on nov.brand =oct.brand and nov.month <> oct.month
12 where nov.sales_per_brand >oct.sales_per_brand
13

```

INFO : Map 1: 5/5 Map 3: 5/5 Reducer 2: 1/1/2 Reducer 4: 2/2
INFO : Completed executing command(queryId=hive_20210206113423_707d41de-1608-430e-967c-57636b19e109), Time taken:
4.517 seconds
INFO : OK

Query History Saved Queries Query Builder Results (100+)

nov.brand	
1	ardell
2	art-visage
3	balbcare
4	beauty-free
5	beautyblender

8. Your company wants to reward the top 10 users of its website with a Golden Customer plan. Write a query to generate a list of top 10 users who spend the most on purchases.

```

select user_id,sum(price) as purchase_Amount from part_events where event_type =
'purchase'
group by user_id order by purchase_Amount desc limit 10;

```

Add a name...
Add a description...

14.93s Database ecom ▼ Type text ▼ ⚙️ ?

```

1 select user_id,sum(price) as purchase_Amount from part_events where event_type = 'purchase'
2 group by user_id order by purchase_Amount desc limit 10;
3

```

```

INFO : Map 1: 5/5      Reducer 2: 2/2 Reducer 3: 1/1
INFO : Map 1: 5/5      Reducer 2: 2/2 Reducer 3: 1/1
INFO : Completed executing command(queryId=hive_20210206105917_1c25d021-6bee-431a-9f59-dd3638b556ad), Time taken:
4.939 seconds
INFO : OK

```

application_1612593054396_0014

Query History
Saved Queries
Query Builder
Results (10)

	user_id	purchase_amount
1	557790271	2715.8699957430363
2	150318419	1645.970008611679
3	562167663	1352.8499938696623
4	531900924	1329.4499949514866
5	557850743	1295.4800310581923
6	522130011	1185.3899966478348
7	561592095	1109.700007289648
8	431050134	1097.5900000333786

Step 13 : Dropped all tables and database and Terminated the cluster. I have tried to do all this in one go so as to prevent heavy costing due to extended use.